

Managing GitHub with Terraform

...

Who am I?

What is Terraform?



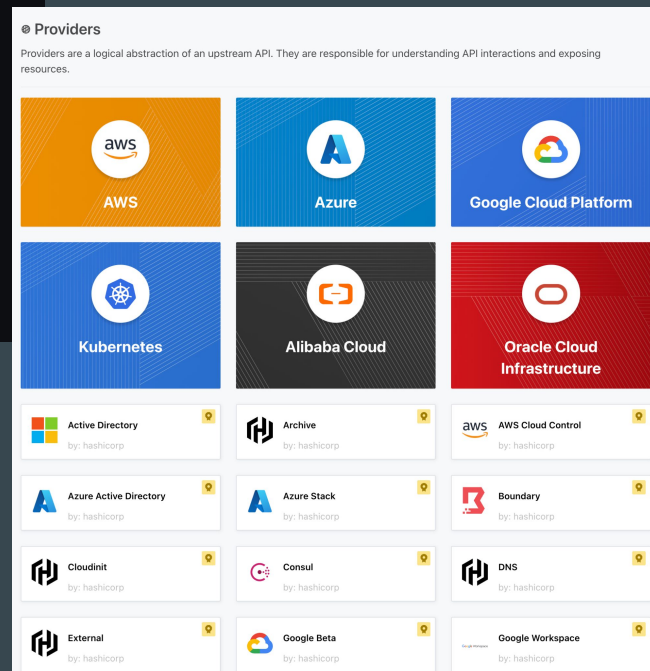
What is Terraform?

Terraform is an infrastructure as code tool that lets you build, change, and version cloud and on-prem resources safely and efficiently.

v1.5.x (latest) ▾

Terraform manages resources through providers
<https://registry.terraform.io/browse/providers>

HashiCorp Terraform is an infrastructure as code tool that lets you define both cloud and on-prem resources in human-readable configuration files that you can version, reuse, and share. You can then use a consistent workflow to provision and manage all of your infrastructure throughout its lifecycle. Terraform can manage low-level components like compute, storage, and networking resources, as well as high-level components like DNS entries and SaaS features.



Terraform Resources

- Resources are the most important element in the Terraform language.
- Each resource block describes one or more infrastructure objects
 - compute instances
 - virtual networks
 - DNS records.

```
resource "aws_instance" "web" {  
    ami           = "ami-a1b2c3d4"  
    instance_type = "t2.micro"  
}
```

Create an ec2 instance

<https://developer.hashicorp.com/terraform/language/resources/syntax>

Terraform State

- All terraform provisioned resources are stored in a state file.
- The state file can be stored in different places, configured as a [backend](#). Some examples include;
 - local file (all in terraform.tfstate locally)
 - S3 backed by dynamodb for state locks
- State files need to be treated as secure as they contain confidential information about your resources.
- As working with terraform requires access to the state file, storing it in a common, secure location is necessary for teams.

Example Configuration

```
terraform {  
  backend "s3" {  
    bucket = "mybucket"  
    key    = "path/to/my/key"  
    region = "us-east-1"  
  }  
}
```

Copy 

bucket = s3 bucket name to store the state

key = path in s3 to the state file

region = the region the bucket exists in

dynamodb_table = The table to store the state locks (missing in the above example), [doc here](#)

Include the Provider

- Terraform providers need to be required..
- Running terraform init will pull the required providers.
 - terraform init also manages provider updates
 - You can specify specific versions using the [version constraints](#)
- It's possible for the provider to accept configuration options via parameters.
 - With GitHub it's possible to override the token with a variable.

```
terraform {  
  required_providers {  
    github = {  
      source = "integrations/github"  
      version = "~> 5.0"  
    }  
  }  
}  
  
# Configure the GitHub Provider  
provider "github" {}
```

```
provider "github" {  
  token = var.token # or `GITHUB_TOKEN`  
}
```

GitHub Provider

...

Some benefits of managing GitHub with Terraform

- Automation and Consistency
 - The Terraform code is stored in GitHub and deployed by conventional CI/CD methods.
 - This consistency helps to maintain a reliable and predictable development environment.
- Simplified Resource Management
 - Terraform's declarative language abstracts the complexities of interacting with GitHub's API.
 - This abstraction simplifies the management of repositories, teams and members
- Versioning and Auditing
 - As Terraform code is version-controlled (i.e. stored in GitHub), it facilitates easy tracking of changes made to GitHub resources.
 - This audit trail allows teams to understand who made changes, when they were made, and why they were made, improving transparency and accountability.
- Scalability and Reusability
 - As infrastructure needs grow, Terraform enables the simple scaling of GitHub resources.
 - Terraform goes further and modules allow the creation of reusable and parameterised configurations which reduces duplication of effort and provides some consistency across projects.
- Safe Changes and Rollbacks
 - Terraform's plan and apply workflow lets you preview changes before applying them to GitHub.
 - This feature allows you to assess potential impacts and catch errors before they affect the environment.
 - In case of issues, you can roll back changes to a known working state easily.
- Infrastructure as Code (IaC) Approach
 - The above points are really points that sum up IaC.
 - With Terraform, GitHub resources are defined as code, allowing for version control, collaboration, and consistency across environments.
 - IaC promotes better documentation, repeatability, and reduces the risk of manual configuration errors.

Getting around

- The [GitHub Provider docs](#)

github

GITHUB DOCUMENTATION

Filter

[github provider](#)

> Resources

> Data Sources

Overview

GitHub Provider

The GitHub provider is used to interact with GitHub resources.

The provider allows you to manage your GitHub organization's members and teams easily. It needs to be configured with the proper credentials before it can be used.

Use the navigation to the left to read about the available resources.

- Documentation is maintained by the provider
- Pretty well maintained

github provider

Resources

- github_actions_environment_secret
- github_actions_environment_variable
- github_actions_organization_oidc_subject_claim_customization_template
- github_actions_organization_permissions
- github_actions_organization_secret
- github_actions_organization_secret_repositories
- github_actions_organization_variable
- github_actions_repository_access_level
- github_actions_repository_oidc_subject_claim_customization_template
- github_actions_repository_permissions
- github_actions_runner_group
- github_actions_secret
- github_actions_variable
- github_app_installation_repositories
- github_app_installation_repository
- github_branch
- github_branch_default
- github_branch_protection
- github_branch_protection_v3
- github_codespaces_organization

github provider

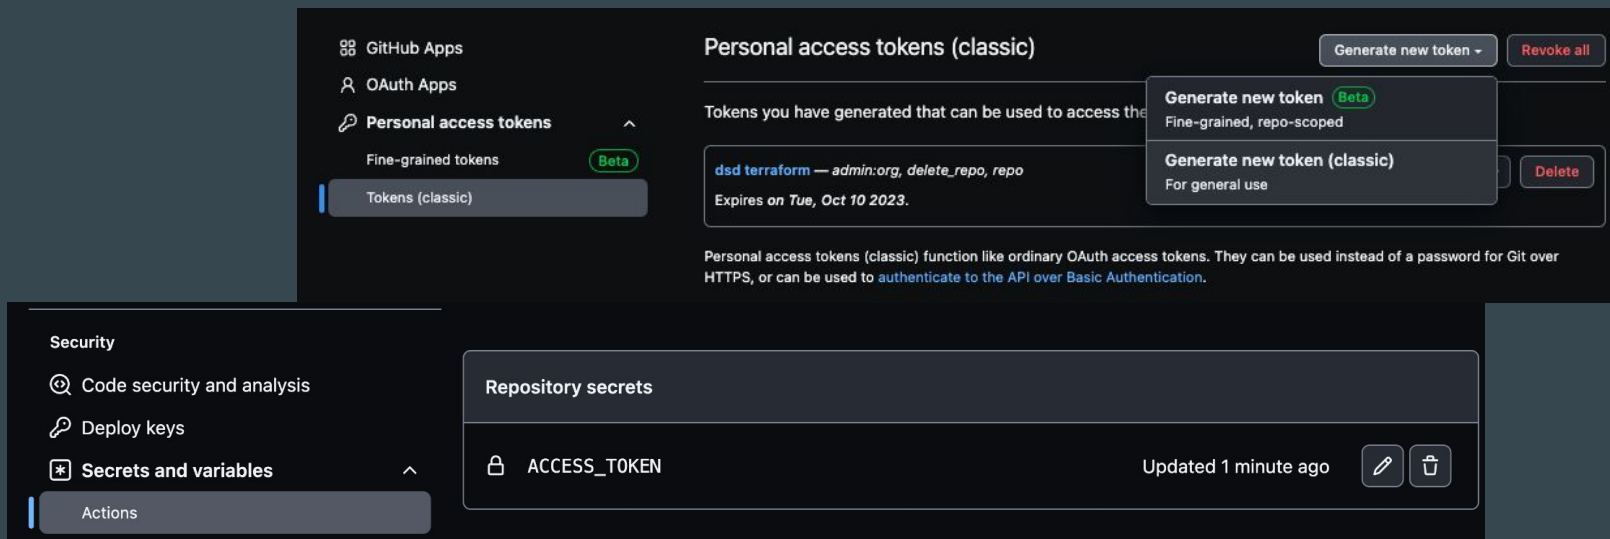
Resources

Data Sources

- github_actions_environment_secrets
- github_actions_environment_variables
- github_actions_organization_oidc_subject_claim_customization_template
- github_actions_organization_public_key
- github_actions_organization_registration_token
- github_actions_organization_secrets
- github_actions_organization_variables
- github_actions_public_key
- github_actions_registration_token
- github_actions_repository_oidc_subject_claim_customization_template
- github_actions_secrets
- github_actions_variables
- github_app
- github_app_token
- github_branch
- github_branch_protection_rules
- github_codespaces_organization_public_key
- github_codespaces_organization_secrets

The Setup

- As a service account;
- Generate a classic token
- Stored the token in the repository secrets in GitHub



Code Segue


Some Enterprise

Benefits of SSO for GitHub

- Centralised Authentication
- Centralised user provisioning
- Immediate user deprovisioning
 - When a user gets removed from AD they lose access to the organisation
- Support for MFA
- Can integrate existing systems like Active Directory

How we integrate with AD

- Well we cheat a little,
- We use OneLogin as our SSO provider
 - OneLogin has a SCIM integration with GitHub
 - This is setup via a [OneLogin application](#) and a GitHub OAuth application
- Our flow for adding a user to our GitHub organisation
 - Add the user to a specific AD group
 - The user then gets and invite to the GitHub organisation.
 - To accept the invite, the user needs to login to OneLogin
 - Select the GitHub tile (now available in the OneLogin UI)
 - They'll be prompted to sign in to GitHub
 - OneLogin stores this link between its user and GitHub



Some things to think about

- We have a manageable number repositories to have in their own files.
How do you scale this further while still being able to manage it? Use a DB?

Q&A